

PDL Lab3:. Binary Classification of Heart Disease of Patients using Deep Neural Network

SWETHA JENIFER 225229142

1. load the dataset

```
In [1]: import pandas as pd
```

```
In [2]: df=pd.read_csv("heart_data.csv")
```

```
In [3]: df.head()
```

Out[3]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [4]: df.shape
```

Out[4]: (303, 14)

```
In [5]: df.size
```

Out[5]: 4242

In [6]: df.info

```
Out[6]: <bound method DataFrame.info of
thalach  exang  oldpeak  \
0      63    1    3      145    233    1      0      150    0      2.3
1      37    1    2      130    250    0      1      187    0      3.5
2      41    0    1      130    204    0      0      172    0      1.4
3      56    1    1      120    236    0      1      178    0      0.8
4      57    0    0      120    354    0      1      163    1      0.6
..      ...    ..    ..      ...    ...    ...    ...    ...    ...    ...
298    57    0    0      140    241    0      1      123    1      0.2
299    45    1    3      110    264    0      1      132    0      1.2
300    68    1    0      144    193    1      1      141    0      3.4
301    57    1    0      130    131    0      1      115    1      1.2
302    57    0    1      130    236    0      0      174    0      0.0

      slope  ca  thal  target
0         0  0    1        1
1         0  0    2        1
2         2  0    2        1
3         2  0    2        1
4         2  0    2        1
..      ...  ..    ..      ...
298        1  0    3        0
299        1  0    3        0
300        1  2    3        0
301        1  1    3        0
302        1  1    2        0

[303 rows x 14 columns]>
```

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

2. Split the dataset

```
In [8]: X=df[['age','sex','cp','trestbps','chol','fbs','restecg','thalach','exang','oldpe
```

```
In [9]: y=df[['target']]
```

```
In [11]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s
```

```
In [12]: X_train.shape
```

```
Out[12]: (242, 13)
```

```
In [13]: X_test.shape
```

```
Out[13]: (61, 13)
```

3. Create a neural network based on the following requirements

```
In [14]: from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense
```

```
In [15]: model = Sequential()  
model.add(Dense(8, input_dim=13, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

4. Compile your model with learning rate = 0.001, optimizer as 'RMSprop', Mean square error loss and metrics as 'accuracy'.

```
In [16]: from tensorflow import keras
```

```
In [17]: optimizer = keras.optimizers.RMSprop(learning_rate=0.001)
```

```
In [18]: model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
         model.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 3s 23ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 2/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 3/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 4/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 5/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 6/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 7/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 8/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 9/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 10/10  
9/9 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy:  
0.5496
```

```
Out[18]: <keras.callbacks.History at 0x22d50f8d270>
```

```
In [19]: model.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.4754 - accuracy:  
0.5246
```

```
Out[19]: [0.4754098355770111, 0.5245901346206665]
```

5. Print the summary of the model: model.summary()

In [20]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 8)	112
dense_1 (Dense)	(None, 1)	9
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		

6. Train the model for 200 epochs and batch size as 10

In [21]: `model.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])`
`model.fit(X_train, y_train, epochs=200, batch_size=10, verbose=1)`

```
Epoch 1/200
25/25 [=====] - 1s 1ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 2/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 3/200
25/25 [=====] - 0s 11ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 4/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 5/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 6/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy: 0.5496
Epoch 7/200
25/25 [=====] - 0s 1ms/step - loss: 0.4504 - accuracy: 0.5496
```

In [22]: `model.evaluate(X_test, y_test)`

```
2/2 [=====] - 0s 4ms/step - loss: 0.4754 - accuracy: 0.5246
```

Out[22]: [0.4754098355770111, 0.5245901346206665]

7. Save the trained model in a variable, such as, history. Also, you can split your training data for validation such as 20% of training data

```
In [24]: history = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=20)

Epoch 1/100
20/20 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 2/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 3/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 4/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 5/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 6/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 7/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
```

8. Evaluate the trained model to predict the probability values for the test data set (ie., xtest and ytest)

```
In [25]: model.evaluate(X_test, y_test)

2/2 [=====] - 0s 2ms/step - loss: 0.4754 - accuracy: 0.5246
```

```
Out[25]: [0.4754098355770111, 0.5245901346206665]
```

9. Print the model accuracy and model loss as below (Use can use the 'history' object we have saved). Sample code is given below.


```
In [26]: history.history.keys()
```


```
Out[26]: dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
In [28]: import matplotlib.pyplot as plt
```

Matplotlib is building the font cache; this may take a moment.

```
In [29]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['Train', 'Validation'])
plt.show()
```





10. Do further experiments

```
In [30]: model1 = Sequential()

model1.add(Dense(16, input_dim=13, activation='relu'))
model1.add(Dense(8, activation='relu'))
model1.add(Dense(1, activation='sigmoid'))
```



```
In [31]: model1.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
         model1.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 2ms/step - loss: 0.5111 - accuracy:  
0.4504  
Epoch 2/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2543 - accuracy:  
0.6198  
Epoch 3/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2492 - accuracy:  
0.6281  
Epoch 4/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2566 - accuracy:  
0.6446  
Epoch 5/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2464 - accuracy:  
0.6240  
Epoch 6/10  
9/9 [=====] - 0s 1ms/step - loss: 0.2396 - accuracy:  
0.6322  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2450 - accuracy:  
0.6198  
Epoch 8/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2513 - accuracy:  
0.6157  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2375 - accuracy:  
0.6488  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.2347 - accuracy:  
0.6446
```

```
Out[31]: <keras.callbacks.History at 0x22d56351240>
```

```
In [32]: model1.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.1707 - accuracy:  
0.7541
```

```
Out[32]: [0.17073306441307068, 0.7540983557701111]
```

In [33]: `history1 = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=20)`

```
Epoch 1/100
20/20 [=====] - 0s 6ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 2/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 3/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 4/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 5/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 6/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
Epoch 7/100
20/20 [=====] - 0s 3ms/step - loss: 0.4560 - accuracy: 0.5440 - val_loss: 0.4286 - val_accuracy: 0.5714
```

In [34]: `model1.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 16)	224
dense_3 (Dense)	(None, 8)	136
dense_4 (Dense)	(None, 1)	9
Total params: 369		
Trainable params: 369		
Non-trainable params: 0		

In [35]: `ls = history1.history`

```
In [36]: new = pd.DataFrame.from_dict(ls)
new
```

Out[36]:

	loss	accuracy	val_loss	val_accuracy
0	0.455959	0.544041	0.428571	0.571429
1	0.455959	0.544041	0.428571	0.571429
2	0.455959	0.544041	0.428571	0.571429
3	0.455959	0.544041	0.428571	0.571429
4	0.455959	0.544041	0.428571	0.571429
...
95	0.455959	0.544041	0.428571	0.571429
96	0.455959	0.544041	0.428571	0.571429
97	0.455959	0.544041	0.428571	0.571429
98	0.455959	0.544041	0.428571	0.571429
99	0.455959	0.544041	0.428571	0.571429

100 rows × 4 columns

```
In [37]: model2 = Sequential()
model2.add(Dense(32, input_dim=13, activation='relu'))
model2.add(Dense(16, activation='relu'))
model2.add(Dense(8, activation='relu'))
model2.add(Dense(1, activation='sigmoid'))
```

```
In [38]: model2.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
model2.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 2/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 3/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5496 - accuracy:  
0.4504  
Epoch 4/10  
9/9 [=====] - 0s 1ms/step - loss: 0.5455 - accuracy:  
0.4545  
Epoch 5/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5455 - accuracy:  
0.4545  
Epoch 6/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5455 - accuracy:  
0.4545  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5494 - accuracy:  
0.4504  
Epoch 8/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5617 - accuracy:  
0.4339  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5502 - accuracy:  
0.4463  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.5454 - accuracy:  
0.4545
```

```
Out[38]: <keras.callbacks.History at 0x22d57637f40>
```

```
In [39]: model2.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.5246 - accuracy:  
0.4754
```

```
Out[39]: [0.5245802402496338, 0.4754098355770111]
```

```
In [40]: model2.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 32)	448
dense_6 (Dense)	(None, 16)	528
dense_7 (Dense)	(None, 8)	136
dense_8 (Dense)	(None, 1)	9

=====
Total params: 1,121

Trainable params: 1,121

Non-trainable params: 0
=====

```
In [41]: model3 = Sequential()  
model3.add(Dense(64, input_dim=13, activation='relu'))  
model3.add(Dense(32, activation='relu'))  
model3.add(Dense(16, activation='relu'))  
model3.add(Dense(8, activation='relu'))  
model3.add(Dense(1, activation='sigmoid'))
```

```
In [42]: model3.compile(loss='mse', optimizer=optimizer, metrics=['accuracy'])  
         model3.fit(X_train, y_train, epochs=10, batch_size=30, verbose=1)
```

```
Epoch 1/10  
9/9 [=====] - 1s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 2/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 3/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 4/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 5/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 6/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 7/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 8/10  
9/9 [=====] - 0s 10ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 9/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496  
Epoch 10/10  
9/9 [=====] - 0s 2ms/step - loss: 0.4504 - accuracy:  
0.5496
```

```
Out[42]: <keras.callbacks.History at 0x22d557c2b90>
```

```
In [43]: model3.evaluate(X_test, y_test)
```

```
2/2 [=====] - 0s 3ms/step - loss: 0.4754 - accuracy:  
0.5246
```

```
Out[43]: [0.4754098355770111, 0.5245901346206665]
```

```
In [44]: model3.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
dense_9 (Dense)	(None, 64)	896
dense_10 (Dense)	(None, 32)	2080
dense_11 (Dense)	(None, 16)	528
dense_12 (Dense)	(None, 8)	136
dense_13 (Dense)	(None, 1)	9
=====		
Total params: 3,649		
Trainable params: 3,649		
Non-trainable params: 0		